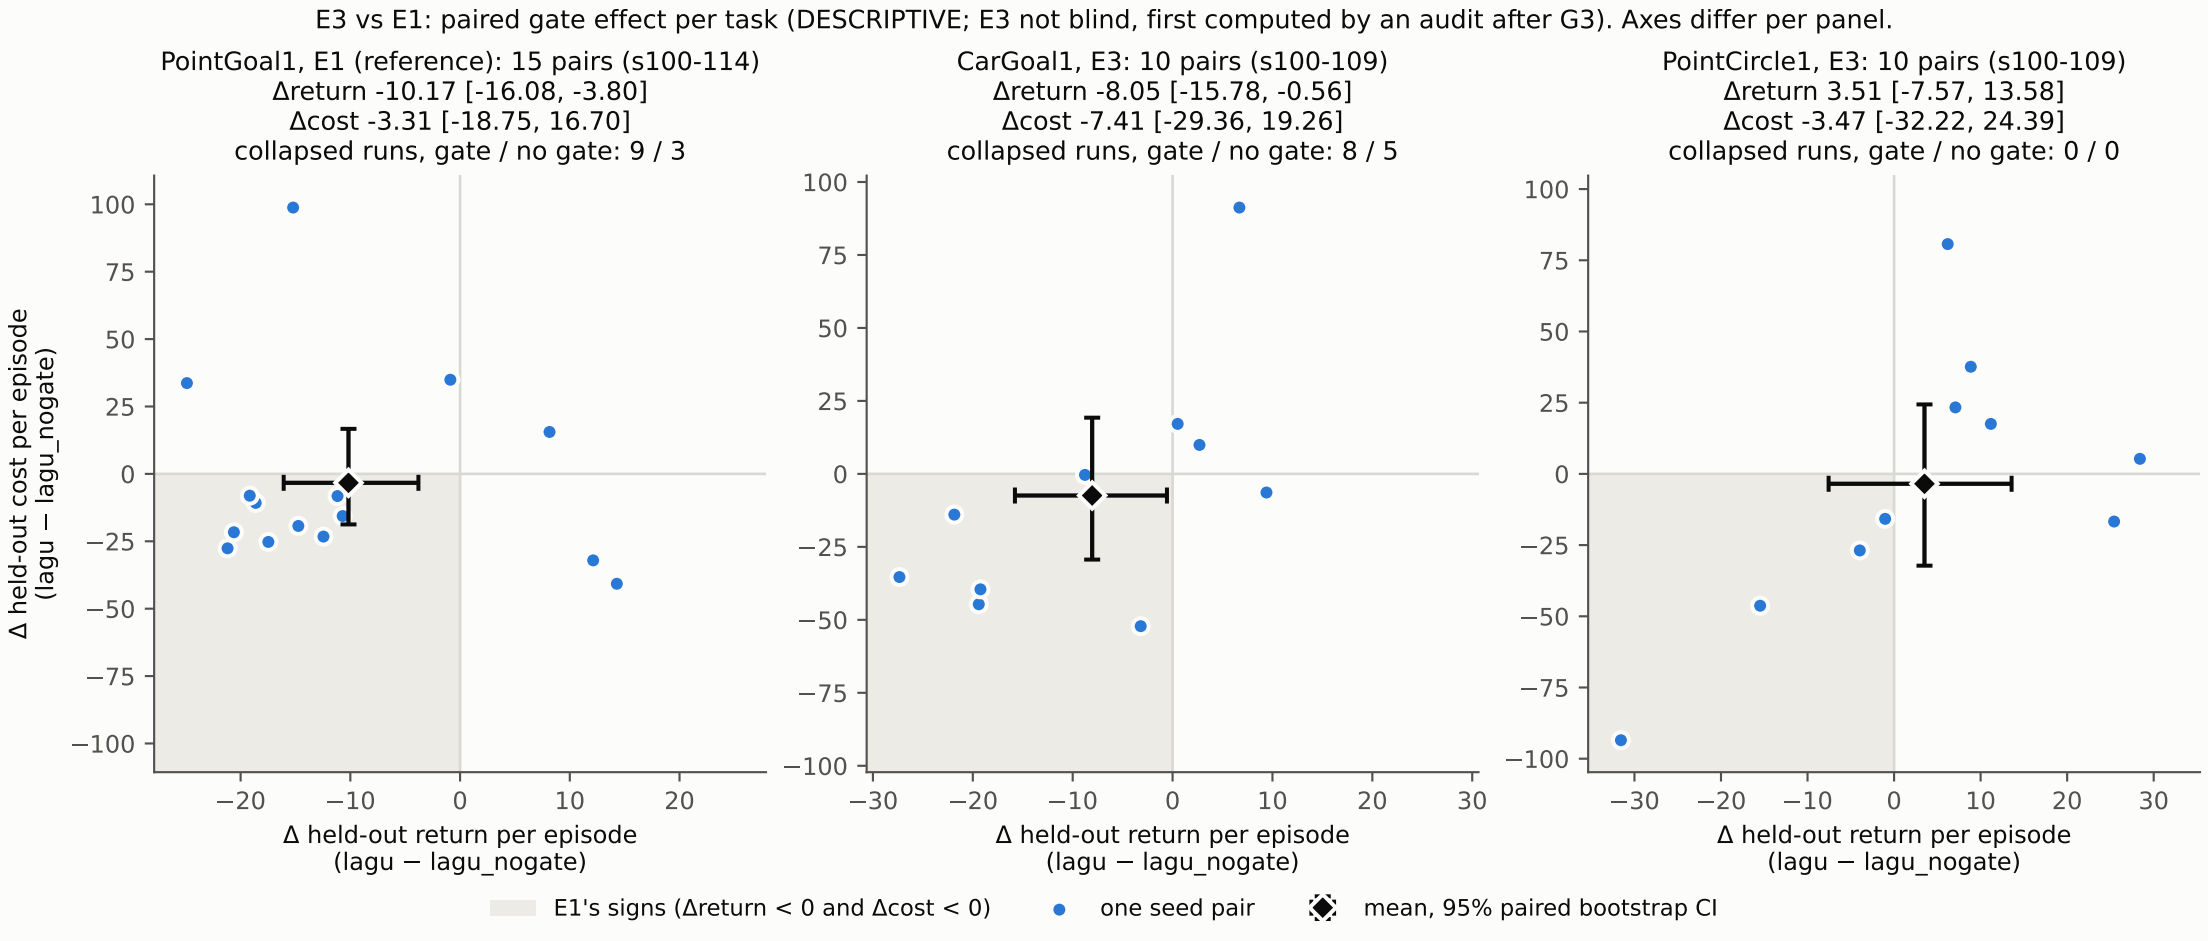
E3 vs E1: paired gate effect per task (DESCRIPTIVE; E3 not blind, first computed by an audit after G3). Axes differ per panel.
PointGoal1, E1 (reference): 15 pairs (s100-114)
Δreturn -10.17 [-16.08, -3.80]
Δcost -3.31 [-18.75, 16.70]
collapsed runs, gate / no gate: 9 / 3
CarGoal1, E3: 10 pairs (s100-109)
Δreturn -8.05 [-15.78, -0.56]
Δcost -7.41 [-29.36, 19.26]
collapsed runs, gate / no gate: 8 / 5
PointCircle1, E3: 10 pairs (s100-109)
Δreturn 3.51 [-7.57, 13.58]
Δcost -3.47 [-32.22, 24.39]
collapsed runs, gate / no gate: 0 / 0
Δ held-out cost per episode (lagu − lagu_nogate)
Δ held-out return per episode (lagu − lagu_nogate)
E1's signs (Δreturn < 0 and Δcost < 0)
one seed pair
mean, 95% paired bootstrap CI