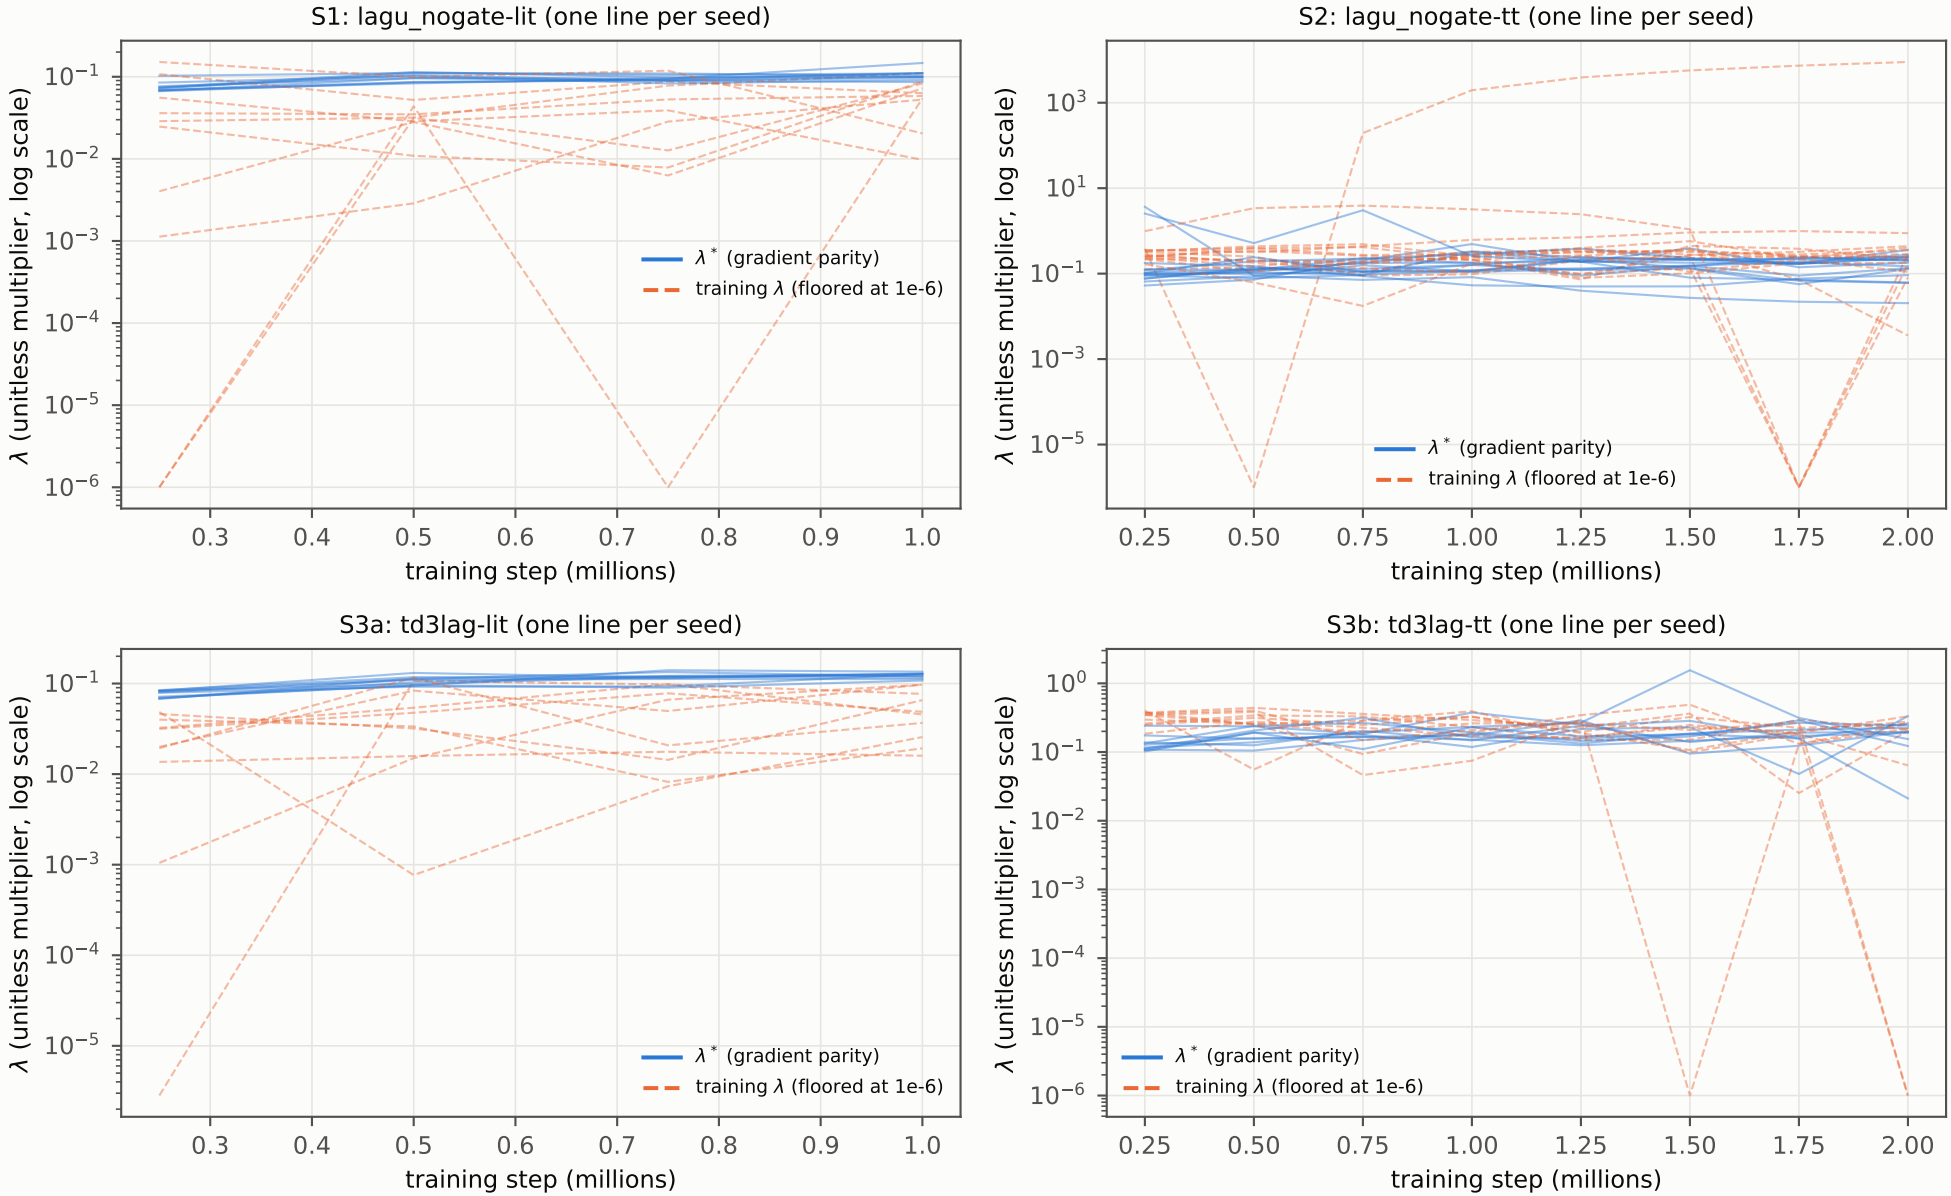
DESCRIPTIVE: gradient parity λ* and the training multiplier over training
S1: lagu_nogate-lit (one line per seed)
S2: lagu_nogate-tt (one line per seed)
S3a: td3lag-lit (one line per seed)
S3b: td3lag-tt (one line per seed)
λ (unitless multiplier, log scale)
training step (millions)
λ* (gradient parity)
training λ (floored at 1e-6)